

Machine Learning Project

Group members:

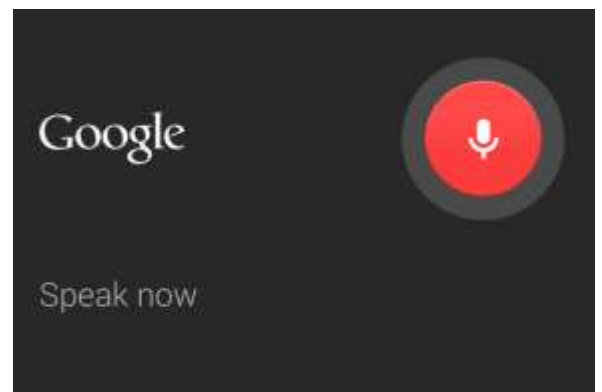
Sidra (EL-015)

Aqsa Sarfaraz(EL-011)

Voice recognizer identify human voice and predict baby male and female voices

Backstory

This branch spreads very fast and now it is almost everywhere – military, education, telephony, in-car systems, security or usual life. The simplest example is our smartphones, almost all smartphones support such thing as speech recognition. I wouldn't say that it is the most important thing and that it is necessary for everyone, I would rather say that it takes place in future technologies, such as full-value AI and security. But this technology is simple only in words. Actually it is difficult enough.



Aim


Generally, speech recognition is wide branch and difficult on the one hand but although very interesting on the another hand. So in this presentation I'd like to show the way we can recognize gender of person's voice.



Dataset

Make a long story short sound is a pressure wave which is created by a vibrating object. Humans can hear sound waves with frequencies between about 20 Hz and 20 kHz. Analyzed frequency range of 0hz-280hz.

The dataset consists of 3168 recorded voice samples, collected from (baby) male and female speakers.




The following acoustic properties of each voice:

- **meanfreq**: mean frequency (in kHz)
- **sd**: standard deviation of frequency
- **median**: median frequency (in kHz)
- **Q25**: first quantile (in kHz)
- **Q75**: third quantile (in kHz)
- **IQR**: interquantile range (in kHz)
- **skew**: skewness (see note in specprop description)
- **kurt**: kurtosis (see note in specprop description)
- **sp.ent**: spectral entropy
- **sfm**: spectral flatness
- **mode**: mode frequency
- **centroid**: frequency centroid (see specprop)
- **peakf**: peak frequency (frequency with highest energy)
- **meanfun**: average of fundamental frequency measured across acoustic signal
- **minfun**: minimum fundamental frequency measured across acoustic signal
- **maxfun**: maximum fundamental frequency measured across acoustic signal
- **meandom**: average of dominant frequency measured across acoustic signal
- **mindom**: minimum of dominant frequency measured across acoustic signal
- **maxdom**: maximum of dominant frequency measured across acoustic signal

- **dfrange**: range of dominant frequency measured across acoustic signal
- **modindx**: modulation index. Calculated as the accumulated absolute difference between adjacent measurements of fundamental frequencies divided by the frequency range
- **label**: (baby) male or female


The most significant features: meanfun and IQR




From all 21 features we can separate “meanfun” and “IQR”. Is it enough of these two features to make predictions? We will see later.

It is possible because of next reasons:

- particular qualities
- accuracy of equipment not enough
- interference during recording
- speech defects



As we could see, there are some samples that belong to male but graph tells us that it is female. And vice versa.



Prediction results

For making prediction I've used 4 algorithms: Logistic Regression, K-means, Naive Bayes and SVM. Also, I found that feature scaling improves results, except K-means.

Below you can see results of each algorithm with 2534 train samples and 634 test samples with all features and with 2 features separately.

Features	Feature scaling	K-means	Logistic Regression
All	Without	54.416 % clusters: 2	88.17 %
All	With	53.31 % clusters: 1	97.003 %
'meanfun', 'IQR'	Without	88.96 % clusters: 2	90.22 %
'meanfun', 'IQR'	With	53.31 % clusters: 1	95.58 %

CODING AND OUTPUTS

```
import pandas as pd

voice = pd.read_csv('voice.csv')
print(voice.head())
voice["label"].value_counts()
```

```
meanfreq      sd      median      Q25      ...      maxdom      dfrange      modindx      label
0  0.059781  0.064241  0.032027  0.015071  ...  0.007812  0.000000  0.000000  male
1  0.066009  0.067310  0.040229  0.019414  ...  0.054688  0.046875  0.052632  male
2  0.077316  0.083829  0.036718  0.008701  ...  0.015625  0.007812  0.046512  male
3  0.151228  0.072111  0.158011  0.096582  ...  0.562500  0.554688  0.247119  male
4  0.135120  0.079146  0.124656  0.078720  ...  5.484375  5.476562  0.208274  male

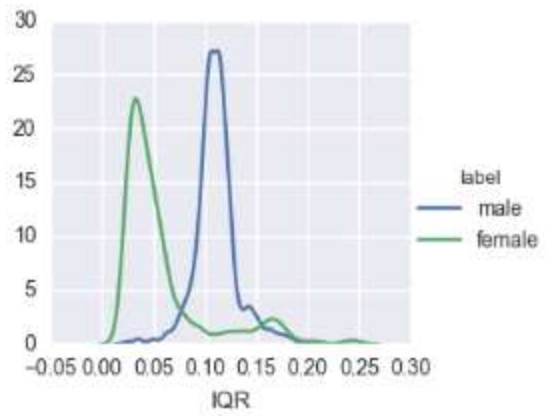


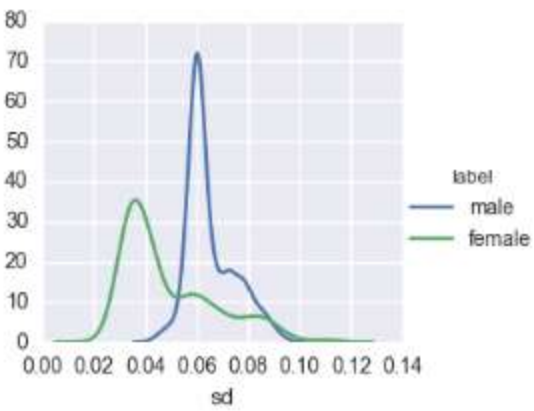


[5 rows x 21 columns]
Out[12]:
male      1584
female    1584
Name: label, dtype: int64
```







```
voice.info()
```









```
In [13]: voice.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3168 entries, 0 to 3167
Data columns (total 21 columns):
meanfreq      3168 non-null float64
sd             3168 non-null float64
median        3168 non-null float64
Q25           3168 non-null float64
Q75           3168 non-null float64
IQR           3168 non-null float64
skew          3168 non-null float64
kurt          3168 non-null float64
sp.ent        3168 non-null float64
sfm           3168 non-null float64
mode          3168 non-null float64
centroid      3168 non-null float64
meanfun       3168 non-null float64
minfun        3168 non-null float64
maxfun        3168 non-null float64
meandom       3168 non-null float64
mindom        3168 non-null float64
maxdom        3168 non-null float64
dfrange       3168 non-null float64
modindx       3168 non-null float64
label         3168 non-null object
dtypes: float64(20), object(1)
memory usage: 519.9+ KB
```

```
import seaborn as sns
from seaborn import plt


for col in voice.columns[:-1]:
    sns.FacetGrid(voice, hue="label", size=3).map(sns.kdeplot, col).add_le
gend()
plt.show()
```









```
sns.FacetGrid(voice, hue="label", size=7).map(plt.scatter, "IQR", "meanfun")\n.add_legend()\nplt.show()
```



```
sns.FacetGrid(voice, hue="label", size=7).map(plt.scatter, "meanfreq", "meanfun").add_legend()
plt.show()
```



```
from sklearn.preprocessing import LabelEncoder

# replace male/female => 1/0
gender_encoder = LabelEncoder()
voice['label'] = gender_encoder.fit_transform(voice.iloc[:, -1])
from sklearn.cross_validation import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.cross_validation import train_test_split

y=voice['label']
x = voice.iloc[:, :-1]
scaler = StandardScaler()
scaler.fit(x)
x = scaler.transform(x)
```

```

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, r
andom_state=2)

# code without scale below

#train, test = train_test_split(voice, test_size=0.2, random_state=2)

# separating data in features and labels
#x_train = train.iloc[:, :-1]
#y_train = train.iloc[:, -1]
#x_test = test.iloc[:, :-1]
#y_test = test.iloc[:, -1]

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

lr = LogisticRegression()
lr.fit(x_train, y_train)
prediction = lr.predict(x_test)

print('LR result: ', accuracy_score(y_test, prediction))

```

```
LR result: 0.970031545741
```

```

from sklearn.cluster import KMeans
import math

result = 0
for num_clusters in range(1, len(x_train)):
    kmeans = KMeans(n_clusters=num_clusters, random_state=0)
    kmeans.fit(x_train, y_train)
    prediction = kmeans.predict(x_test)
    cur_result = accuracy_score(y_test, prediction)
    if cur_result < result or math.fabs(cur_result-result) < 0.01:
        break
    result = cur_result

print('K-means result: ', result, '. Cluster quantity = ', num_clusters-1)

```

```
K-means result: 0.533123028391 . Cluster quantity = 1
```

```

from sklearn.naive_bayes import GaussianNB
nb = GaussianNB()
nb.fit(x_train, y_train)
prediction = nb.predict(x_test)

print('NB result: ', accuracy_score(y_test, prediction))

```

NB result: 0.862776025237

```
from sklearn.svm import SVC

svc=SVC()
svc.fit(x_train,y_train)
prediction = svc.predict(x_test)
print('SVM result: ', accuracy_score(y_test,prediction))
```

SVM result: 0.97476340694

```
voice_cut = voice[['IQR', 'meanfun', 'label']]

#y_cut=voice['label']
#x = voice_cut.iloc[:, :-1]
#scaler = StandardScaler()
#scaler.fit(x)
#x_cut = scaler.transform(x)

#x_train, x_test, y_train, y_test = train_test_split(x_cut, y_cut, test_size=0.2, random_state=2)

train, test = train_test_split(voice_cut, test_size=0.2, random_state=2)

# separating data in features and labels
x_train = train.iloc[:, :-1]
y_train = train.iloc[:, -1]
x_test = test.iloc[:, :-1]
y_test = test.iloc[:, -1]

lr = LogisticRegression()
lr.fit(x_train, y_train)
prediction = lr.predict(x_test)

print('LR result with 2 features: ', accuracy_score(y_test, prediction))
```

LR result with 2 features: 0.902208201893

```
result = 0
for num_clusters in range(1, len(x_train)):
    kmeans = KMeans(n_clusters=num_clusters, random_state=0)
    kmeans.fit(x_train, y_train)
    prediction = kmeans.predict(x_test)
    cur_result = accuracy_score(y_test, prediction)
    if cur_result < result or math.fabs(cur_result-result) < 0.01:
        break
```

```
result = cur_result

print('K-means result with 2 features: ', result, '. Cluster quantity = ',
num_clusters-1)
```

```
K-means result with 2 features:  0.889589905363 . Cluster quantity =  2
```

```
nb = GaussianNB()
nb.fit(x_train, y_train)
prediction = nb.predict(x_test)

print('NB result with 2 features: ', accuracy_score(y_test, prediction))
```

```
NB result with 2 features:  0.958990536278
```

```
svc=SVC()
svc.fit(x_train,y_train)
prediction = svc.predict(x_test)

print('SVM result with 2 features: ', accuracy_score(y_test,prediction))
```

```
SVM result with 2 features:  0.900630914826
```

Conclusion

So, best result results we have in SVM, then logistic regression, Naive Bayes and K-means.

But in case with only 2 features('meanfun', 'IQR') without feature scaling results are more precise. Naive Bayes, for example, gives 95 % against 85-85% with all features. That is because algorithm is based on suggestion of independence. The most significant improvement is in K-means without feature scaling, it grows up from ~54% to 88%. The reason is the same, less quantity of features.

Generally, results of partial dataset don't lose to full, but sometimes gives even better results. Hence, in some cases we can use only 2 features, instead of all 21.